We thank reviewers R1, R2, R3 for their substantial time investment and for their incisive feedback.

Limits. R1 highlights ways our precision must improve. Thm2 states: ***{{****For each d, every non-deg. local min.* $\theta_\star$ *has a nbhd U whose every member* θ_0 *induces, via Thm1, a T-indexed sequence of degree-d truncations* $f_T \in \mathbb{R}[\eta]$ *that converges as* $T \to \infty$ *to some* $f_d \in \mathbb{R}[\eta]$*.**}}*** $\mathbb{R}[\eta]$ is the free polynomial ring in $\dim(\mathcal{M})^2$ many variables, topologized as euclidean space. So if $L_{d,T}(\eta)$ is Thm1's truncation, Thm2 controls $L_d(\eta) = \lim_{\tilde{T}\to\infty} L_{d,\tilde{T}}(\eta)$ but not $L_T(\eta) = \lim_{\tilde{d}\to\infty} L_{\tilde{d},T}(\eta)$, even for $d, T \gg 1$. Thm1,2's significance stems from **empirical** findings that their *formal* power series [**Wi**] bear on SGD practice (w.r.t. which *any* infinities are idealizations). Our mathematics was but a strong heuristic,[1] so we didn't examine when L_d, L_T agree. Still: **Prop A**. ***{{****Fix* $U \subseteq \mathcal{M}$ *open,* $\theta_\star \in U$ *a non-deg. local min. of l. Assume §B.1 as well as global, prob.-1 bounds* $(|l_x(\theta)|, \|\nabla l_x(\theta_\star)\|) < C$. ***If*** *some* $Q_-, Q_+ \in SPD$ *bound the hessian* $(Q_- < \nabla\nabla l_x(\theta) < Q_+)$ *on U,* ***then*** $\forall d$ *and for any* θ_0 *in some nbhd* V_d *of* $\theta_\star$*:* $\exists T_0, g$ *with* $|g(T)|$ *in* $\exp(-big\Omega(T))$ *so that* $\sup_{T\geq T_0} |L_d(\eta) - L_T(\eta) - g(T)|$ *(exists on some nbhd in SPSD of* $\eta = 0$ *and) is* $o(\eta^d)$*.**}}*** **SP(S)D** consists of symmetric positive (semi)definites.

Sharp Minima. Like us, R3 finds Cor5[2] counterintuitive. SGD's noise consists not of weight *displacements* but of error terms $\nabla l_x(\theta) - \nabla l(x)$ in the *gradient* estimate; compare Fig5  to [**Ke**]'s Fig1. Say θ's 1D with $l(\theta) = a\theta^2/2$ and training loss $\hat{l}(\theta) = l(\theta) + b\theta$. At $\hat{l}$'s min. $\theta = -b/a$, $l(\theta) = b^2/(2a)$. So for fixed b, sharp min.a ($a \gg 1$) overfit less (demo here). The covariance C controls b^2, explaining Cor5's $C/2H$ factor. Here, optimization *to convergence* favors sharp min.a ($\star$); convergence is slow at flat min.a, so flat min.a also overfit little ($\diamond$). (Our small-η assumption rules out the possibility that H is so sharp that SGD diverges: we treat $\eta H \ll 1$). Prior work (Pg12Par5, e.g. [**Ke**] and [**Di**]) supports both pro-flat and pro-sharp intuitions. Recognizing η's role in translating gradients to displacements, we account for both ($\star$) and ($\diamond$) and hence unify existing intuitions (§4.3). We view it as a merit that our theory makes such counterintuitive phenomena visible.

Comparison to ODE. We thank R2 for [**Ba**]. [**Ba**]'s LemA.3 specializes our LemKey. In our language, [**Ba**] computes fuzzless diagrams (noiseless ≡ cumulants vanish ≡ fuzz-having diagrams vanish). The only such η^2 diagrams for $l(\theta_T)$ are , . The N^{-1} correction to [**Ba**]'s noiseless assumption consists of diagrams with one fuzzy outline (Pg24Tab1).

Notation. R3 recognizes our expectands as tensor expressions; they are often fully contracted (so scalar) and are always random variables in some $\mathbb{R}^k$. Per R2,R3, we'll disemploy 'Einstein notation' and cite [**Cu**] (+ a new §D) for tensor examples. If advised, we'll also forgo diagrams: e.g. $[a][ab : c : d][bcd]$ for  (letters name edges). R3, Pg6Thm2 defines '**non-degenerate**' as '$H \in$ SPD'.

Organization. R2,R3 stress the paper's narrative challenge. We'll arrange the paper into 3 self-contained tracks, each pertinent to a different goal: TrkA [pgs 1-4], for casual readers, will eschew diagrams, theorems, and §1.1/§2.2's heavy notations; illustrate Taylor series via §2.1's proof; identify §3.3's terms; state Cor4 (w/ §B.1's assumptions explicit, w/ PrpA's precision); explain §4.2's curl effect. TrkB [pgs 1-4, 5-12], for seekers of physical intuition, will use TrkA to motivate (and §A.4 to illustrate) §1.1/§2's definitions; relegate §2.2/2.1's LemKey/discussion to §B; add to §2.3.1 a resummation cartoon à la Fig5,7. For space, §C'll absorb §4. TrkC [pgs 5-12, 15-45], for our theory's extenders, will include PrpA (per R1) and more explicit statements and arguments throughout.

References. [**Ba**] D.G.Barrett, B.Dherin. *Implicit Gradient Regularization*. ICLR 2021. [**Cu**] P.McCullagh. *Tensor Methods in Statistics*, §1.1-1.4,§1.8. Dover 2017. [**Di**] L.Dinh, R.Pascanu, S.Bengio, Y.Bengio. *Sharp Minima Can Generalize for Deep Nets*, §1,§5. ICML 2017. [**Ke**] N.S.Keskar et alia. *Large-Batch Training for Deep Learning*, §4. ICLR 2017. [**Wi**] H.Wilf. *Generatingfunctionology*, §2.1-2.3. Academic Press 1994.

1. By clcal.mech. (CM) of thermal continua, ice cubes have energy=∞ [McQuarrie '97, §1-1]. But CM gives real insight.
2. i.e.: *that* ***overfitting*** ($\triangleq l(\theta_T) - l(\theta_0)$ where θ_0 is a min. of l) *has an* η^2 *term greatest when* ηH *has moderate eigenvalues.*